


An algorithm for virtual machine consolidation

Ilya Pautner

Higher School of Economics,
Faculty of Computer Science,
Applied math and informatics


Scientific advisor:
Andrey Tikhonov,
Lead engineer,
Russian Research Institute,
Huawei

Relevance

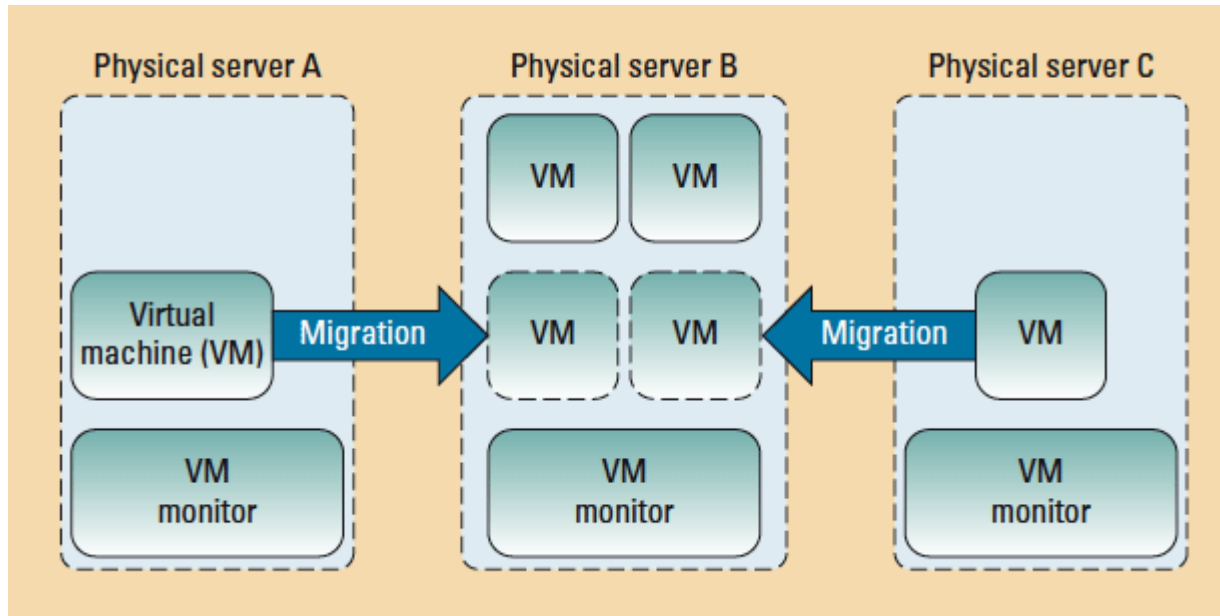


Vector bin packing problem (well-known NP-hard)

- VMs (demands) W : $\vec{w}_i = (w_{1,i}, \dots, w_{d,i})$
- Hosts (capacities) H : $\vec{c}_j = (c_{1,j}, \dots, c_{d,j})$
- Mapping $f: W \rightarrow H$
- Constraint $\forall j \in H, \forall i \in f^{-1}(j): \sum w_i \leq c_j$
- Active score $\#\{f^{-1}(j) \neq \emptyset\}$
- Minimize active score



Virtual machine consolidation (not so well-known)



- Items already are placed in bins (Init_Mapping)
- Migration score: total memory of moved VMs

Mathematical problem statement

Indexing VMs as v_i ($i = 1, \dots, n$) and PMs as p_j ($j = 1, \dots, m$), the following binary variables are introduced:

$$Alloc_{i,j} = \begin{cases} 1 & \text{if } v_i \text{ is allocated on } p_j \\ 0 & \text{otherwise} \end{cases}$$

$$Active_j = \begin{cases} 1 & \text{if } p_j \text{ is active} \\ 0 & \text{otherwise} \end{cases}$$

$$Migr_i = \begin{cases} 1 & \text{if } v_i \text{ is migrated} \\ 0 & \text{otherwise} \end{cases}$$

Using these variables, the integer program can be formulated as follows ($i = 1, \dots, n$ and $j = 1, \dots, m$):

$$\min \quad \alpha \cdot \sum_{j=1}^m Active_j + \mu \cdot \sum_{i=1}^n Migr_i \quad (4) \rightarrow \text{Optimization objective}$$

$$\text{s. t.} \quad \sum_{j=1}^m Alloc_{i,j} = 1 \quad \forall i \quad (5) \rightarrow \text{Virtual machine should be on exactly host}$$

$$Alloc_{i,j} \leq Active_j \quad \forall i, j \quad (6) \rightarrow \text{Host with any VMs should be active}$$


$$\sum_{i=1}^n load(v_i) \cdot Alloc_{i,j} \leq_d cap(p_j) \quad \forall j \quad (7) \rightarrow \text{Capacity constraint}$$

$$Migr_i = 1 - Alloc_{i, map_0(v_i)} \quad \forall i \quad (8) \rightarrow \text{Migration is when VM is moved from one host to another}$$

$$Alloc_{i,j}, Active_j, Migr_i \in \{0, 1\} \quad \forall i, j \quad (10) \rightarrow \text{Integrality constraint}$$

Existing solutions

ILP Solver Search Procedure




19

ILP solver

Well-developed industry

Slow

Branch-and-Bound




Branch and bound

Uses domain knowledge

Slow


Each node in branch-and-bound is a new MIP



First-Fit-Decreasing heuristic

Simple

High migration cost



SerCon heuristic


Smallest migration cost

Ineffective packing

Solution in thesis

Two stages


- Bin packing via FFD
- Migration optimization



Migration types

Internal migration – source host would be active

External migration – source host would NOT be active



Algorithms in comparison

- Initial (do nothing)
- First-Fit-Decreasing
- First-Fit-Decreasing + Migration Optimizer
- VPSolver
- VPSolver + Migration Optimizer
- SerCon

Algorithms in comparison

- Initial (do nothing)
- First-Fit-Decreasing
- First-Fit-Decreasing + Migration Optimizer
- VPSolver
- VPSolver + Migration Optimizer
- SerCon

Q: Why no SerCon + Migration Optimizer?

A: Migration optimizer can reduce internal migrations, and SerCon produces only external migrations.

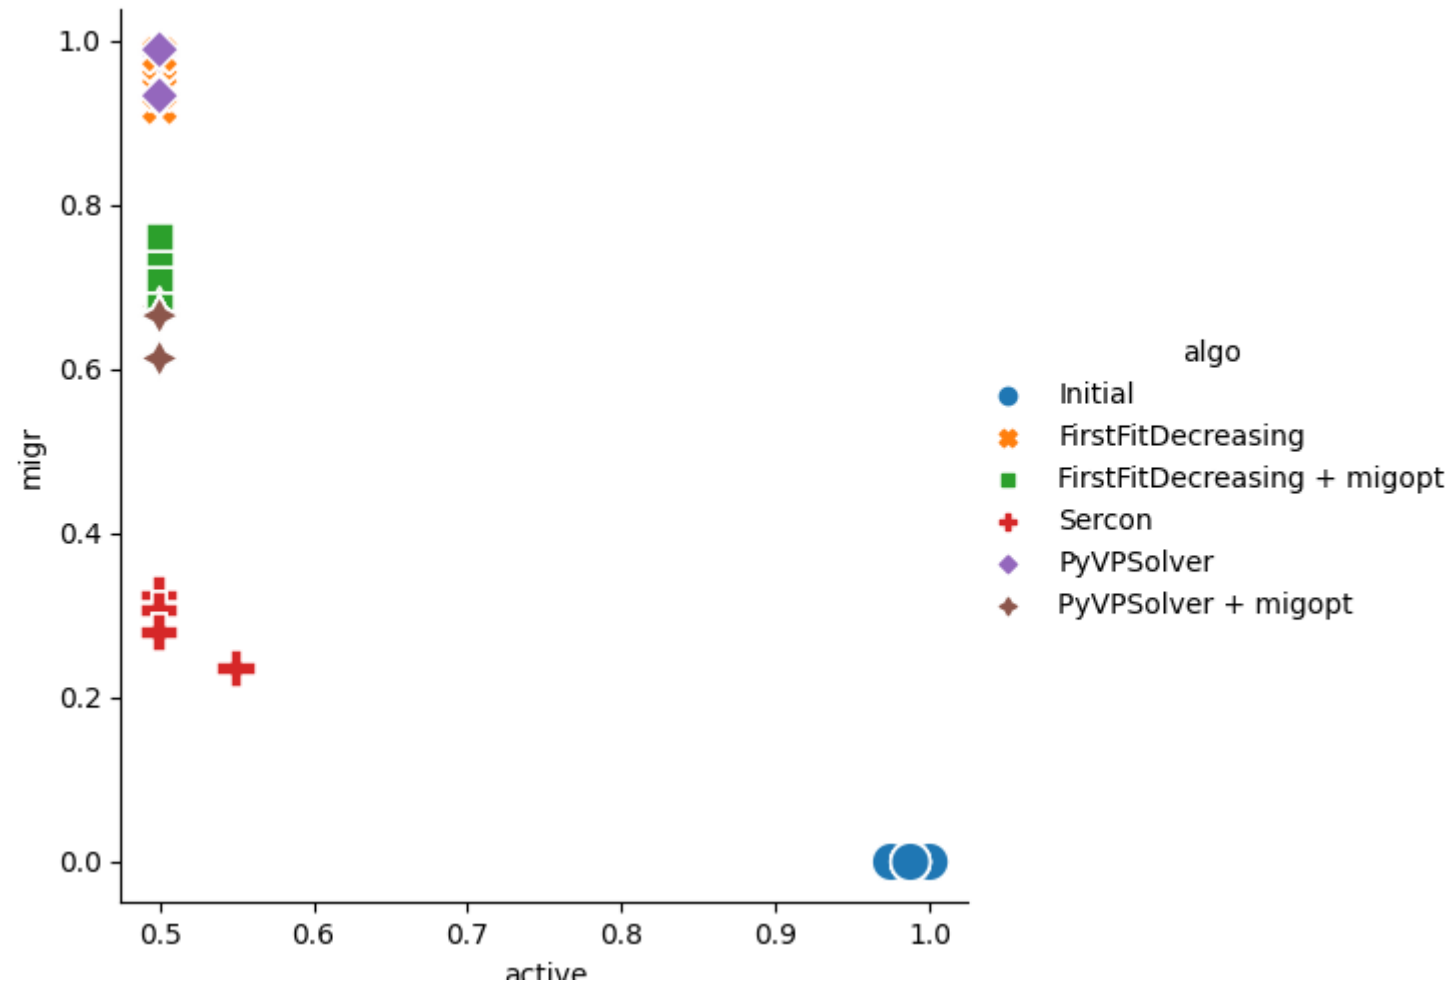
Test types

- Shrink test type
 - Fill hosts with random virtual machines from distribution until first fail
 - Reduce each virtual machine requests by random factor
- Full-pack test type
 - Fill hosts with random virtual machines from distribution until full

Two variants of distributions of virtual machine sizes.

They are based upon real-world statistics of large cloud provider.

Results (on 1 test)



Future research


- Different scores:
 - fixed VM amount score: how many VMs on certain size can be put into this configuration.
 - balancing score: how far the load distribution across the hosts from uniform.
- Better first step heuristic than FFD.
- Migration optimizer improvements.

Main points

- Formulated combinatorial optimization problem
- Provided application for cloud systems
- Reduced to assignment problem
- Pareto curve in results
- Demonstrated 30% decrease in migration score


Migration optimizer ideas (questions section)

- Mapping can be arbitrarily changed if active hosts remain the same.
- Hole system (akin to partition of unity in topology):
One can partition host to non-overlapping parts of smaller size.
- Each hole can accommodate VMs of size not bigger than itself.



Hole making rules


- VM of size n produces hole of size n
- Free space is divided in holes of largest VM size + remainder



Initial and final hole system

- Initial hole system is done from initial mapping.
- Final hole system is done from final mapping, using only active hosts.
- Mapping change between initial and final mapping is a matching between holes in initial and final hole systems: which VM goes where.
- Since we have weights assigned to migrations, it is natural to assign weights to the edges.
 - infinity, if source hole does not fit on destination hole
 - 0, if source hole and destination hole are on same host
 - memory, if source hole and destination hole are on different hosts

Min-weight maximum matching aka Assignment problem



Optimal matching in red.

Total weight:

$$0 + 0 + 4 = 4$$